

# Neural Networks: Representation

5 questions

---

1  
point

1.


Which of the following statements are true? Check all that apply.

- ☐ A two layer (one input layer, one output layer; no hidden layer) neural network can represent the XOR function.
  - ☐ The activation values of the hidden units in a neural network, with the sigmoid activation function applied at every layer, are always in the range (0, 1).
  - ☐ Any logical function over binary-valued (0 or 1) inputs  $x_1$  and  $x_2$  can be (approximately) represented using some neural network.
  - ☐ Suppose you have a multi-class classification problem with three classes, trained with a 3 layer network. Let  $a_1^{(3)} = (h_\Theta(x))_1$  be the activation of the first output unit, and similarly  $a_2^{(3)} = (h_\Theta(x))_2$  and  $a_3^{(3)} = (h_\Theta(x))_3$ . Then for any input  $x$ , it must be the case that  $a_1^{(3)} + a_2^{(3)} + a_3^{(3)} = 1$ .
- 

1  
point

2.

Consider the following neural network which takes two binary-valued inputs  $x_1, x_2 \in \{0, 1\}$  and outputs  $h_{\Theta}(x)$ . Which of the following logical functions does it (approximately) compute?




- ☐ NAND (meaning "NOT AND")
- ☐ AND
- ☐ OR
- ☐ XOR (exclusive OR)

---

1  
point

3.

Consider the neural network given below. Which of the following equations correctly computes the activation  $a_1^{(3)}$ ? Note:  $g(z)$  is the sigmoid activation function.




- ☐  $a_1^{(3)} = g(\Theta_{1,0}^{(2)} a_0^{(2)} + \Theta_{1,1}^{(2)} a_1^{(2)} + \Theta_{1,2}^{(2)} a_2^{(2)})$
- ☐  $a_1^{(3)} = g(\Theta_{1,0}^{(2)} a_0^{(1)} + \Theta_{1,1}^{(2)} a_1^{(1)} + \Theta_{1,2}^{(2)} a_2^{(1)})$
- ☐  $a_1^{(3)} = g(\Theta_{1,0}^{(1)} a_0^{(2)} + \Theta_{1,1}^{(1)} a_1^{(2)} + \Theta_{1,2}^{(1)} a_2^{(2)})$
- ☐  $a_1^{(3)} = g(\Theta_{2,0}^{(2)} a_0^{(2)} + \Theta_{2,1}^{(2)} a_1^{(2)} + \Theta_{2,2}^{(2)} a_2^{(2)})$

1  
point

4.

You have the following neural network:



You'd like to compute the activations of the hidden layer  $a^{(2)} \in \mathbb{R}^3$ . One way to do so is the following Octave code:

```
% Theta1 is Theta with superscript "(1)" from lecture
% ie, the matrix of parameters for the mapping from layer 1 (input) to layer 2
% Theta1 has size 3x3
% Assume 'sigmoid' is a built-in function to compute 1 / (1 + exp(-z))

a2 = zeros (3, 1);
for i = 1:3
    for j = 1:3
        a2(i) = a2(i) + x(j) * Theta1(i, j);
    end
    a2(i) = sigmoid (a2(i));
end
```

You want to have a vectorized implementation of this (i.e., one that does not use for loops). Which of the following implementations correctly compute  $a^{(2)}$ ? Check all that apply.

- ☐  $z = \text{Theta1} * x$ ;  $a2 = \text{sigmoid}(z)$ ;
- ☐  $a2 = \text{sigmoid}(x * \text{Theta1})$ ;
- ☐  $a2 = \text{sigmoid}(\text{Theta2} * x)$ ;

☐  $z = \text{sigmoid}(x)$ ;  $a_2 = \text{sigmoid}(\Theta_1 * z)$ ;

---

1  
point

5.

You are using the neural network pictured below and have learned the


parameters  $\Theta^{(1)} = \begin{bmatrix} 1 & 0.5 & 1.9 \\ 1 & 1.2 & 2.7 \end{bmatrix}$  (used to compute  $a^{(2)}$ ) and

$\Theta^{(2)} = \begin{bmatrix} 1 & -0.2 & -1.7 \end{bmatrix}$  (used to compute  $a^{(3)}$ ) as a function of  $a^{(2)}$ .

Suppose you swap the parameters for the first hidden layer between its two

units so  $\Theta^{(1)} = \begin{bmatrix} 1 & 1.2 & 2.7 \\ 1 & 0.5 & 1.9 \end{bmatrix}$  and also swap the output layer so

$\Theta^{(2)} = \begin{bmatrix} 1 & -1.7 & -0.2 \end{bmatrix}$ . How will this change the value of the output  $h_{\Theta}(x)$ ?



- ☐ It will stay the same.
  - ☐ It will increase.
  - ☐ It will decrease
  - ☐ Insufficient information to tell: it may increase or decrease.
- 



I understand that submitting work that isn't my own may result in permanent failure of this course or deactivation of my Coursera account. Learn more about Coursera's Honor Code